Tool Developer Qualification Course

Reverse Engineering project

Bomb

Joshua Abraham

Taught by by
Dr. Joseph Santmeyer

November 18, 2017

Contents

1 Summary

I began my analysis by reading the assembly of each stage to get a general idea of what they were doing. This static analysis is my preferred way to solve challenges like these. Some of the assembly code snippets in this document has been edited to be more readable (i.e. identifying local variables with a name like "input" instead of an rbp offset).

```
1  nm bomb > symbols.out
2  gdb −batch −ex 'file ./bomb.patched' −ex 'disassemble /r stage_1' > objectdumps/stage_1.out
3  ....
4  gdb −batch −ex 'file ./bomb.patched' −ex 'disassemble /r stage_8' > objectdumps/stage_8.out
```

While performing dynamic analysis with the GNU Debugger (gdb), the program pretends to segfault by detecting that it is being debugged. The program uses ptrace to do so, a common anti-reverse engineering tactic. While the majority of my analysis of the code was static, just reading the code and translating to higher level constructs, I occasionally needed to perform tests with the binary. To avoid having to manually step over the call to ptrace in _start every single time, I patched the binary with gdb to call main *if it was being debugged*, instead of if it was not. I did so with the following commands:

```
1   cp bomb bomb.patched
2   gdb −write −q ./bomb.patched
3   disassemble /r _start
4   ....
5   0x0000000000401461 <+49>:      75 60     jne     0x4014c3
6   ....
7   set {unsigned char}0x0000000000401461 = 0x74
8   disassemble /r _start
9   ....
10  0x0000000000401461 <+49>:      74 60     je      0x4014c3
11  ....
12  quit
```

Using an x64 assembler I found that the je and jne instructions are the same size, and only differ by one byte. We can skip the fake segfault message by making a patched copy of the bomb binary that will run in a debugger (note that the patched binary will only work in a debugger).

My solutions to the stages are as follows:

1. Stage 1: swordfish

2. Stage 2: jabraham

3. Stage 3: 1872

4. Stage 4: 107 214 428 856 1713

5. Stage 5:   1172

6. Stage 6: 48 a 48

7. Stage 7: (press enter)

8. Stage 8: (run ./stag8.sh in the same directory as the binary, then press enter).

9. Stage ?: Did not complete.

2 Stage 1 Analysis

Analysis:

Stage 1 is straightforward. Below is an abbreviated listing of the relevant disassembled code:

```
1 mov      esi , 0x4018f8     ; "swordfish"
2 mov      rdi , rax          ; user input
3 call     strcmp
4 test     eax , eax
5 setz     al
6 leave
```

This stage calls strcmp with user input and the string "swordfish". If the return value is zero (meaning the strings are identical), this stage returns non-zero and the program proceeds.

Solve Script:

Since there was no math needed to calculate my password for this stage, no script was needed.

3   Stage 2 Analysis

Analysis:

Stage 2 starts with a call to getenv. This function takes one parameter and returns the value of a specified environment variable.

```
1  mov     edi, 0x401902      ; "USER"
2  call    getenv
3  mov     QWORD PTR [rip+0x201264], rax
```

The user name is stored in a global variable that is used throughout the binary at rip+0x201264. In my case the user name is "jabraham". This and the saved user input are passed to strcmp, similarly to Stage 1.

```
1  mov     rax, [rbp-0x8]
2  mov     rsi, rdx           ; "jabraham"
3  mov     rdi, rax           ; "jabraham"
4  call    strcmp
5  test    eax, eax
6  setz    al
7  leave
```

After the call to strcmp, the program checks the return value in eax and sets al if the zero flag is set. The stage is passed if the user supplies their own username.

Solve Script:

As with Stage 1, there was no need to perform calculations to find the password for this stage.

4 Stage 3 Analysis

Analysis:

Stage 3 takes user input and saves it in [rbp-0x28]. The global variable holding the username is moved into [rbp-0x8].

```
1 push    rbp
2 mov     rbp, rsp
3 sub     rsp, 30h
4 mov     [rbp-0x28], rdi
5 mov     rax, fs:28h
6 mov     [rbp-0x08], rax
7 xor     eax, eax
8 mov     rcx, cs:src
9 mov     rax, [rbp-0x28]
```

Once these variables are set, they are passed to the C function strncat. It takes three parameters: the destination, the source, and the maximum number of bytes to concatenate.

```
1 mov     edx, 0x80          ; number
2 mov     rsi, rcx           ; source (username)
3 mov     rdi, rax           ; destination (input)
4 call    strncat
```

The destination will now look similar to this: "100jabraham". The bomb then calls strlen on the original username variable.

```
1 mov     rax, [rbp+0x28]
2 mov     rdi, rax           ; username ("jabraham")
3 call    strlen
```

The length of the username is saved on the stack at [rbp-0x10]. Then the bomb gets ready for a call to sscanf. This function scans a string using a format string and stores each match in following parameters. Here sscanf is passed the concatenated string, the format specifier "%u", and the address of a stack variable. The "%u" tells us that sscanf is looking for a number, meaning the password for this stage is a single number.

```
1 mov     [rbp-0x10], rax
2 mov     [rbp-0x14], 0
3 lea     rdx, [rbp-0x14]
4 mov     rax, [rbp-0x28]
5 mov     esi, 0x401907   ; "%u"
6 mov     rdi, rax
7 mov     eax, 0
8 call    sscanf
```

Next the binary performs series of calculations with the total length and input number. Note: for readability, I have changed the offsets to reflect each stack variable's purpose.

```
1  mov     eax, [rbp+input_as_a_number]
2  mov     eax, eax
3  mov     edx, 0
4  div     [rbp+input_len]
5  mov     [rbp+input_as_a_number], eax
6  mov     eax, [rbp+input_as_a_number]
7  shr     eax, 2
8  mov     [rbp+input_as_a_number], eax
9  mov     eax, [rbp+input_as_a_number]
10 mov     edx, 0AAAAAAABh
11 mul     edx
12 mov     eax, edx
13 shr     eax, 1
14 cmp     rax, [rbp+input_len]
15 setnbe  al
```

Let us decode what is happening in this block of code. The input number is moved to eax. Then it is divided by the total length of the concatenated string. The result is stored in the input number variable, and a logical shift right 2 is performed on it. This value is then multiplied by the value 0xAAAAAAAB and shifted right 1. The calculated value is compared to the original input length and al is set to one if it is greater. To summarize the previous in a more readable format, in order to pass this stage your input must satisfy the following:

$$\frac{input\_number \div input\_len}{12} > input\_len$$

If so, the function returns non zero and the bomb continues to stage 6.

Solve Script:

In order to solve the stage I wrote a script in Python 3 to quickly perform the math needed. The script is called "stage3.py" and is located in this repository for your convenience.

```python
from os import getenv
username = getenv("USER")          # Calculate the length of the username
usernamelen = len(username)
ans = 0                            # Start at zero

# Loop until the first number is greater than
# the length of the number plus the username length

while True:
    ans += 1
    length = len(str(ans)) + usernamelen
    if (ans // length // 12) > length:
        break

print(ans)
```

On the UMBC server, this code gives me the integer value 1872.

5 Stage 4 Analysis

Analysis:

Stage 4's disassembly has conditionals and jumps, making a little bit more difficult to read but still straightforward.

```
1 mov     eax,DWORD PTR ds:0x201170        ; the username
2 movzx   eax, byte ptr [rax]
3 movsx   eax, al
4 mov     [rbp+first_letter_in_username], eax
```

stage_4 starts by saving the first character in the username in a local stack variable.

```
1 mov     r9, rsi
2 mov     r8, rcx
3 mov     rcx, rax
4 mov     rdx, r10
5 mov     esi, 0x40190a     ; "%u %u %u %u %u"
6 mov     eax, 0
7 call    sscanf
8 add     rsp, 10h
9 cmp     eax, 5
10 je     0x9
```

The binary then invokes the sscanf function again, this time with the format string "%u %u %u %u %u". This means that we need 5 integers to pass the stage. In fact, the binary checks if the function returns 5 and fails if not.

```
1 cmp     [rbp+counter], 4
2 jbe     0x401023
3 ...
4 mov     rax, [rbp+counter]
5 mov     eax, [rbp+rax*4+input_numbers]
6 cmp     eax, [rbp+first_letter_in_username]
7 jbe     0x40103d
8 ...
9 mov     rax, [rbp+counter]
10 mov    eax, [rbp+rax*4+input_numbers]
11 add    [rbp+first_letter_in_username], eax
12 jmp    0x401044
13 ...
14 add    [rbp+counter], 1   ; continue back at top
```

This code shows the path of successful execution of stage_4. This can be thought of as a for loop that executes five times, over each number in our input. During each pass through the loop, the value is added to the ASCII value of the first character in the user name. If the next value is smaller than the previous, the stage fails.

```
1 mov     rdi, rax          ; username
2 call    strlen
3 mov     rcx, rax
4 mov     rax, rbx
5 mov     edx, 0
6 div     rcx
7 mov     rax, rdx
8 test    rax, rax
9 setz    al
```

After the loop, the stage calls strlen on the username and then divides the last sum by the length. If the remainder (in rdx) is 0, the stage returns non-zero. The pseudocode below shows an approximation of what is happening here:

Pseudocode:

```
1  for number in input_numbers:
2    if number <= first_letter:
3      return False
4    else:
5      first_letter += number
6  return first_letter % strlen(username) == 0
```

Solve Script:

I wrote the following script to solve this challenge:

```
1  from os import getenv
2  # Get the length of the username
3  username = getenv("USER")
4  length = len(username)
5  first_letter = ord(username[0])
6
7  nums = [first_letter + 1]
8
9  for i in range(4):
10     nums.append(nums[-1] + (nums[-1] ))
11
12 # Add one if the last element is odd
13 if (nums[-1] % 2 == 0):
14     nums[-1] += 1
15
16 print(" ".join(str(x) for x in nums))
```

This script gives me the following output on the server: 107 214 428 856 1713

6 Stage 5 Analysis

Analysis:

Stage 5 initially uses calls used many times in earlier stages, so I will skip over the minutiae. It calls sscanf on the user input to retrieve three values with the following format string: "%c %d %c". The stage expects a letter, followed by an integer, and another character. There is a call to strlen to get the length of the username.

```
1 mov      rax , [ rbp+index ]
2 cmp      rax , [ rbp+username_len ]
3 jb       0x40110a
```

This is the main loop of the stage and it terminates when the counter is not smaller than the username length.

```
1 mov      rax , QWORD PTR [ rax∗8+0x401928 ]
2 jmp      rax
```

This took a bit of research but these sources put me in the right direction:

- https://stackoverflow.com/questions/3012011/switch-case-assembly-level-code

- https://stackoverflow.com/questions/9815448/jmp-instruction-hex-code

The rest of the code represents a jump table based on the letters in the username. I look through the binary to find the offsets, and it is in .rodata. You can use the objdump tool to examine sections, so I run the following command:

```
1 objdump −s −j  . rodata bomb.patched > objectdumps /. rodata
```

I write a python script where I annotate each jump for each value. It is called stage5.py. Once the code gets to the offset, one of the input characters is added to or subtracted from the input number. The calculation depends on the ASCII value of each character in the username. The stage passes if the final value is zero. The actual calculations are in the stage5.py file.

Solve Script:

```python
1 from os import getenv
2 username = getenv("USER")    # Get username
3 username_vals = [ ord( character ) for character in username ]
4 input_num = 1000              # Just pick a high number. We will adjust later
5 master_input = 1000
6 char1 = ord ( '~' )                    # Use tildes because it 's the highest printable ascii val (126)
7 char2 = ord ( '~' )
8
9 def math_one ( val ) :
10     global input_num
11     input_num −= ( char1 + val )
12
13 def math_two ( val ) :
14     global input_num
15     input_num −= ( val ∗ 2)
16
17 def math_three ( val ) :
18     global input_num
19     input_num −= ( val )
20
21 def math_four ( val ) :
22     global input_num
23     input_num += char2 − val
24
```

```
25  def math_five(val):
26      global input_num
27      input_num -= (val + char2)
28
29  def math_six(val):
30      global input_num
31      input_num -= (char2)
32
33  group_one = [97,101,105,111,117,121,math_one]
34  group_two = [98,99,100,math_two]
35  group_three = [102,103,104,math_three]
36  group_four = [106,107,108,109,110,math_four]
37  group_five = [112,113,114,115,116,math_five]
38  group_six = [118,119,120,122,math_six]
39
40  master_list = [group_one, group_two, group_three, group_four, group_five, group_six]
41
42  for value in username_vals:
43      group = [x for x in master_list if value in x]       #find our list
44      group[0][-1](value)
45
46  ans = master_input - input_num
47  print("{} {} {}".format(chr(char1), ans, chr(char2)))
```

This script gives me " 1172 " on the server.

7 Stage 6 Analysis

Analysis:

Stage 6 begins with a call to sscanf on the user input. The equivalent C code is as follows:

```
sscanf(in_string, "%d %c %d", num1, character, num2);
```

We can see that we are expected to pass an integer, a character, and an integer to the stage. Next, stage_6 calls a function called ___ with the two integers from user input.

```
mov     edx, [rbp+num_two]
mov     eax, [rbp+num_one]
mov     esi, edx
mov     edi, eax
call    ___
```

This function is recursive and makes several calls back to itself. Looking through the objdumb we can see the termination criteria:

```
cmp     [rbp+first_number], 0
jnz     0x1a
...
cmp     [rbp+second_number], 0
jnz     0x14
...
mov     rax, QWORD PTR [rip+0x200eb8]       ; username
movzx   eax, byte ptr [rax]
movsx   eax, al
...
ret
```

This was my first time seeing recursion in assembly, but the key is to look for what causes the recursion to walk back up to the original caller. In this case, we can see that happens if the parameters are 0. The ___ function then returns the first character in the username, in this case 'j'. Now, back in stage_4:

```
mov     [rbp+first_letter], eax
movzx   eax, [rbp+input_character]
movsx   eax, al
and     eax, [rbp+first_letter]
mov     ecx, [rbp+num_one]
mov     edx, [rbp+num_two]
add     edx, ecx
cmp     eax, edx
setz    al
```

The return value from ___ is put in a local variable, and a does a bitwise and with it and the input character. If this value equals the sum of the two input integers, the stage is passed. The pseudocode looks like this:

```
return (first_letter & input_character) == num_one + num_two
```

Since we have 'j' (106) 106 logical and & the character must equal the sum of our numbers lets pick a number that we know & 106 that will be even (to get equal factors) will land us in the printable range (48 - 126):

$$106 \& 103 == 98$$
$$98/2 = 49$$
$$103 == g$$

So a working solution is "49 g 49". I wrote a Python 3 script to automate the math.

Solve Script:

The script starts with the ASCII value of 'a', and increments up until a match is found.

```python
from os import getenv

user_char = getenv("USER")[0]          # Get first char in username
user_char = ord(user_char)

startchar = ord('a')

while True:
    if (user_char & startchar) % 2 == 0:
        char = (user_char & startchar) // 2
        print("{} {} {}".format(char, chr(startchar), char))
        break
    else:
        startchar += 1
```

This gives me "48 a 48" on the server.

8   Stage 7 Analysis

Analysis:

Stage 7 required me to carefully read the disassembly in order to find the solution. Unlike previous stages, the is no sscanf to signal the format that the stage is expecting. They key in this stage is to work backwards from the return. We know that stage_7 needs to return non-zero in order to pass the stage, so we look for any time the eax register is manipulated before the ret.

```
1 cmp       [rbp-0x8], rax
2 setz      al
```

Here is one such instruction at 0x40142b. There are two ways to get to this code, but one of them is interesting:

```
1 mov       rax, [rbp+input]
2 movzx     eax, byte ptr [rax]
3 test      al, al
4 jnz       0x4013c7
```

If the input is zero, the jump is not taken, allowing us to fall through to the set instruction. The user input isn't manipulated anywhere else, so if there is no input the stage passes.

Solve Script:

None was needed for this stage.

9 Stage 8 Analysis

Analysis:

Stage 8 was a challenge that required dynamic analysis. Reading ahead in the disassembly, we can see several "call syscall" instructions. Instead of using sscanf to get user input, this stage appears to be using system calls provided by the kernel to get the user password. I used gdb to complete this stage.

```
gdb −q ./bomb.patched
(gdb) b stage_8
...
<give each stage the correct flag here>
...
(gdb) si <until 0x4016fe>
(gdb) disas
(gdb) x/s 0x401ab0
0x401ab0:        "%d %c %d"
```

We saw stage_8 call malloc and check the return value. Now we are at the first call to syscall, and we see that it is getting 4 parameters: 0x2, "%d %c %d", and 0x40000. Looking up the Linux x64 syscall table (https://filippo.io/linux-syscall-table/) we can see that 0x2 is the "open" call. This call opens files, so we can guess that solving this stage requires a file. There is a trick here, the stage moves 0x6020e0 to a local variable. Running "x/s 0x6020e0" in gdb shows us that this is "stage_8_input.txt". However, the Linux x64 Application Binary Interface requires 6 registers to be used before passing stack parameters to a function call, so this is a red herring. I run the following command to create the required file:

```
touch %d\ %c\ %d
```

Stepping over the system call, we see that if the open is successful, we jump to a setup for another system call. Run "disas" in gdb to see the parameters for this call. We see 0, the return value from open, a variable from the earlier malloc, and a number (0x80). Using the system call reference, we can see that this is a read call. Data read from the file descriptor from open will be put into the memory from the malloc. So the binary is reading from the "%d %c %d" file. Looking at the objdump output, we can see that there is a check on the return value as follows:

```
0x0000000000401757 <+148>:mov      QWORD PTR [rbp−0x8],rax
0x000000000040175b <+152>:cmp      QWORD PTR [rbp−0x8],0x3f
0x0000000000401760 <+157>:jg       0x401778 <stage_8+181>
```

If the return value is greater than 0x3f (63_{10}), we continue. I run the following command to put 80 bytes of text in the file:

```
python2 −c "print 'x' * 80" >> %d\ %c\ %d
```

Back in gdb, step over the read and through the check. We are now at this block of code:

```
0x0000000000401778 <+181>:mov      rax,QWORD PTR [rbp−0x10]
0x000000000040177c <+185>:mov      rsi,rax
0x000000000040177f <+188>:mov      edi,0x3
0x0000000000401784 <+193>:mov      eax,0x0
0x0000000000401789 <+198>:call     0x400900 <syscall@plt>
0x000000000040178e <+203>:mov      QWORD PTR [rbp−0x10],rax
0x0000000000401792 <+207>:cmp      QWORD PTR [rbp−0x10],0x0
```

We can see one more syscall with the following parameters: 3 and the file descriptor. System call 3 is close, and this command closes the file. Step through this. We then see a value moved into rbp-0x20. In gdb use "x/6c $rbp-0x20" to view the characters ('5$,1$3'). After stepping in gdb, we can see that the first byte from the file is compared with 0x41 and if they are equal, the stage fails. This check will not fail unless the file had 'A' in it. Continue stepping and we see that the stage does a logical xor of the byte from magic string with 0x41 and compares with the byte from the file:

```
1  0x00000000004017d2  <+271>:mov     rax ,QWORD PTR  [rbp−0x20]
2  0x00000000004017d6  <+275>:movzx   eax ,BYTE PTR  [rax]
3  0x00000000004017d9  <+278>:xor     eax ,0x41
4  0x00000000004017dc  <+281<:cmp     dl ,al
5  0x00000000004017de  <+283>:je      0x4017b6  <stage_8+243>
```

Now we know that the binary is decoding the magic string with a simple xor, so we can find the correct input by xor-ing each byte of the $rbp-0x20 variable:

$$'5' \oplus 0x41 =' t'$$
$$'\$' \oplus 0x41 =' e'$$
$$',' \oplus 0x41 =' m'$$
$$'1' \oplus 0x41 =' p'$$
$$'\$' \oplus 0x41 =' e'$$
$$'3' \oplus 0x41 =' r'$$

So our input file must have the word "temper". Let's start over but with a correct input file:

```
1  echo temper > %d\ %c\ %d
2  python2 −c "print 'x' * 80" >> %d\ %c\ %d
3  gdb ./bomb.patched −q
4  b 0x00000000004017d9
5  r
6  ...
7  <give correct flags here>
8  ...
9  <step until the string check is done>
```

Remembering past lessons, we need stage_8 to return non-zero. We can see the goal at address 0x00000000004017f6 (mov eax, 0x1). In order for that instruction to be reached, the jump prior can't be taken. Let's take a look at the code:

```
1  0x00000000004017eb  <+296>:mov     rax ,QWORD PTR  [rbp−0x20]
2  0x00000000004017ef  <+300>:movzx   eax ,BYTE PTR  [rax]
3  0x00000000004017f2  <+303>:test    al ,al
4  0x00000000004017f4  <+305>:jne     0x4017fd  <stage_8+314>
5  0x00000000004017f6  <+307>:mov     eax ,0x1
```

If the test al, al sets the Zero Flag, the jump does not happen. This means eax has to have a zero in it. In order for that to happen, the input file needs a null byte after the magic string. Putting together everything we know, this script is what we need:

Solve Script:

```
1  echo temper > %d\ %c\ %d                        # Create and put the xor's hiddenmessage
2  printf "\0" >> %d\ %c\ %d                        # Append a null bute
3  python2 −c "print 'a' * 80" >> %d\ %c\ %d        # to meet the 63 byte requirement
```

10   Secret Stage Analysis

Analysis:

Unfortunately I did not have enough time to tackle the secret stage. My analysis is that the stage performs checks on global variable referenced throughout the bomb and in order to reach the stage, you must supply specially crafted input to earlier stages.

11 Works Cited

References

[1] Switch case assembly level code. (n.d.). Retrieved November 18, 2017, from https://stackoverflow.com/questions/3012011/switch-case-assembly-level-code .

[2] JMP instruction - Hex code. (n.d.). Retrieved November 18, 2017, from https://stackoverflow.com/questions/9815448/jmp-instruction-hex-code